

Cassandra

A few very practical examples

Create and verify key spaces

AT NODE #0:

```
CREATE KEYSPACE one_replicated
WITH REPLICATION = {'class' : 'SimpleStrategy',
                    'replication_factor' : 1};

CREATE KEYSPACE two_replicated
WITH REPLICATION = {'class' : 'SimpleStrategy',
                    'replication_factor' : 2};

SELECT * FROM system.schema_keyspaces;
```

Prepare first table

AT NODE #0:

```
CREATE TABLE one_replicated.users (  
  id int,  
  name text,  
  surname text,  
  PRIMARY KEY(id, surname))  
WITH CLUSTERING ORDER BY (surname ASC)  
AND read_repair_chance = 0.0;  
  
INSERT INTO one_replicated.users(id, name, surname)  
VALUES (0, 'Jan', 'Kowalski');
```

Where does Kowalski live?

AT NODE #0 WITH #1 AND THEN #2 DISABLED; THEN AT #1 WITH @0 DISABLED:

CONSISTENCY ONE;

SELECT * FROM users;

SELECT * FROM users WHERE id = 0;

Where does Kowalski live?

AT NODE #0 WITH #1 AND THEN #2 DISABLED; THEN AT #1 WITH @0 DISABLED:

CONSISTENCY ONE;

SELECT * FROM users; #will fail

SELECT * FROM users WHERE id = 0; #may fail
(too strong CL or RPC timeout – FD not fast enough)

Kowalski has a neighbour!

AT NODE #0

```
INSERT INTO one_replicated.users(id, name, surname)  
VALUES (1, 'Tomasz', 'Iksiński');
```

```
SELECT * FROM users WHERE id = 0  
ORDER BY surname DESC;
```

Kowalski has a neighbour!

AT NODE #0

```
INSERT INTO one_replicated.users(id, name, surname)
VALUES (1, 'Tomasz', 'Iksiński');
```

```
SELECT * FROM users ORDER BY surname DESC;
```

#why can't this just work?!

Kowalski has a neighbour!

AT NODE #0

```
INSERT INTO one_replicated.users(id, name, surname)
VALUES (1, 'Tomasz', 'Iksiński');
```

```
SELECT * FROM users WHERE id = 0 ORDER BY surname
DESC; #it's OK now
```


Replication

For now Replication Factor (RF) = 2

Let's have some strings

AT NODE #0

```
USE two_replicated;
```

```
CREATE TABLE strings (  
    string text PRIMARY KEY)  
WITH read_repair_chance = 0.1;  
  
CONSISTENCY one;
```

Let's have some strings

AT NODE #0

```
INSERT INTO strings(string) VALUES ('test1');  
INSERT INTO strings(string) VALUES ('test2');  
INSERT INTO strings(string) VALUES ('test3');  
INSERT INTO strings(string) VALUES ('test4');  
INSERT INTO strings(string) VALUES ('test5');
```

And where does 'test5' live?

AT NODE #0

```
SELECT token(string) FROM strings  
WHERE string = 'test5';
```

```
$ nodetool getendopints two_replicated strings 'test5'
```

Consistency

A bit more complicated stuff

A map now

AT NODE #0

```
USE two_replicated;
```

```
CREATE TABLE map (  
    key int PRIMARY KEY,  
    value text)
```

```
WITH read_repair_chance = 0.0; #sorry
```

```
CONSISTENCY ONE;
```

```
INSERT INTO map (key, value) values (0, 'version1'); #and  
where exactly this entry lives?
```

We're updating our map

AT NODE #2 AFTER IT ACKNOWLEDGED FAILURE OF NODE #0

USE two_replicated;

CONSISTENCY ONE;

UPDATE map SET value = 'version2' WHERE key = 0;

SELECT value FROM map WHERE key = 0;

Failure scenarios

AT NODE #2 AFTER IT ACKNOWLEDGED RETURN OF NODE #0 AND FAILURE OF #1

SELECT value FROM map WHERE key = 0; #now only
NODE0 should be queried

#is the result correct?!

How about disabling hinted-handoff?

...by editing config file and restarting
the service

Soft reset

AT NODE #2 WHEN ALL THREE NODES ARE UP & RUNNING

CONSISTENCY ALL;

UPDATE map SET value = 'version1' WHERE key = 0;

SELECT value FROM map WHERE key = 0;

CONSISTENCY ONE;

Soft reset

AT #0, #1 AND #2

SELECT value FROM map WHERE key = 0;# should
produce the same result everywhere

We're updating our map

AT NODE #2 AFTER IT ACKNOWLEDGED FAILURE OF NODE #0 (#1 IS UP)

```
UPDATE map SET value = 'version2' WHERE key = 0;
```

```
SELECT value FROM map WHERE key = 0;# result should  
be up to date
```

Failure scenarios

AT NODE #2 AFTER IT ACKNOWLEDGED RETURN OF NODE #0 AND FAILURE OF #1

SELECT value FROM map WHERE key = 0; #now only
NODE0 should be queried

#the result is obsolete – that's our goal!

Now, how to repair this?!

AT NODE #2 AFTER IT ACKNOWLEDGED RETURN OF NODE #1 (#0 IS STILL UP)

CONSISTENCY QUORUM;

SELECT value FROM map WHERE key = 0; #produces an error as long as #1 is not UP

ALTER TABLE map WITH read_repair_chance = 1.0;

SELECT value FROM map WHERE key = 0; #a few times;
#0 should be finally updated

Let's check if NODE0 has been repaired

AT NODE #2 AFTER IT ACKNOWLEDGED FAILURE OF NODE #1

CONSISTENCY ONE;

SELECT value FROM map WHERE key = 0; #finally the correct result!

Thank you

<http://sirius.cs.put.poznan.pl/~inf89719/cassandra/>